

# Process Programming

## Operating System Lab Spring 2015

Roya Razmnoush

CE & IT Department, Amirkabir University of Technology



# Process Control




- fork and vfork
- wait and waitpid
- exec
- system

# Process Control

## fork,vfork



# fork



# fork (cont.)

- `int fork();`
- The **only way** a new process is created by the Unix kernel.
  - The new process created is called the child process.
- The child is a **copy** of the parent.
  - The child gets a copy of the parent's data space, heap and stack.
  - The parent and child don't share these portions of memory.

# fork (cont.)

- This function is called **once**, but return **twice**.
  - The process ID of the new child (to the parent).
    - A process can have more than one child.
  - 0 (to the child).

# fork Sample

```
main()
{
    int pid;
    pid = fork();
    if (pid < 0)
        // error
    else if (pid == 0)
        //child
    else
        //parent
}
```



# fork (cont.)

- We never know if the child starts executing before the parent or vice versa.
  - This depends on the **scheduling algorithm** used by the kernel.

# vfork

- `int = vfork();`
- It has the same calling sequence and same return values as `fork`.
- The child **doesn't copy** the parent data space.
  - The child runs in the address space of the parent.
- With `vfork`, child runs first, then parent runs.

# Process Control

## wait,waitpid



# exit

- Normal termination
  - Executing a `return` from the main function.
  - Calling the `exit` function.
  - Calling the `_exit` function.
- Abnormal termination
  - Calling `abort`.
  - Receives certain signals.

# exit (cont.)

- `int exit (int state);`
- Sometimes we want the terminating process to be able to notify its parent how it terminated.
- For the `exit` and `_exit` function this is done by passing an `exit status` as the argument to these two functions.
- The parent of the process can obtain the `termination status` from either the `wait` or `waitpid` function.

# Termination Conditions

- Parent terminate before the child
  - The `init` process becomes the parent process of any process whose parent terminated.

# Termination Conditions

- Child terminate before the parent
  - The child is completely disappeared, but the parent wouldn't be able to fetch its termination status.
  - The kernel has to keep a certain amount of information for every terminating process.
  - The process that has terminated, but whose parent has not waited for it, is called **zombie**.

# wait

- When a process terminates, the parent is notified by the kernel sending the parent the **SIGCHLD** signal.
- The parent of the process can obtain the termination status from either the **wait** or **waitpid** function.



# wait (cont.)

- The process that calls wait or waitpid can:
  - Block (if all of its children are still running)
  - Return immediately with the termination status of a child (if a child has terminated)
  - Return immediately with an error (if it doesn't have any child process)

# wait and waitpid

- `int wait (int *statloc);`
- `int waitpid (int pid, int *statloc, int options);`
  - If `statloc` is not a null pointer, the termination status of the terminated process is stored in this location.
- The difference between these two function:
  - `wait` can block, while `waitpid` has an option that prevents it from blocking.
  - `waitpid` doesn't wait for the first child to terminate (it can control which process it waits for)

# Process Control

## exec



# exec

- A process cause **another program** to be executed by calling one of the exec functions.
- When a process calls one of the exec functions, that process is completely **replaced** by the new program.
- The process ID doesn't change across an exec.

# exec functions

- `int execl(char *path, char *arg0, ... /*(char *) 0 */);`
- `int execlp(char *path, char *arg0, ... /*(char *) 0, char *envp[] */);`
- `int execlp(char *filename, char *arg0, ... /*(char *) 0 */);`
- `int execv(char *pathname, char *argv0[]);`
- `int execve(char *pathname, char *argv0[], char *envp[]);`
- `int execvp(char *filename, char *envp[]);`

# Process Control system



# System


- execute a shell command.
- `#include <stdlib.h>`
- `int system(const char *string);`
- `System()` executes a command specified in `string` by calling `/bin/sh -c string`, and returns after the command has been completed.
- Return value: the value returned is -1 on error (e.g. fork failed), and the return status of the command otherwise.

# What is the difference between system & exec?

- `system()` calls out to `sh` to handle your command line, so you can get wildcard expansion, etc. `exec()` and its friends replace the current process image with a new process image.
- With `system()`, your program continues running and you get back some status about the external command you called. With `exec()`, your process is obliterated.
- The `exec` function replace the currently running process image when successful, no child is created (unless you do that yourself previously with `fork()`). The `system()` function does `fork` a child process and returns when the command supplied is finished executing or an error occurs.



# Process Relationship



# Inter Process Communication (IPC)




- shared memory
- mapped memory
- pipe
- fifo
- socket
- message passing
- shared semaphore

# Inter Process Communication (IPC)


## Shared memory



# Shared memory



# Shared memory (cont.)



# Shared memory (cont.)

- `int shmget (key_t key, int size, int flag);`
- A shared memory segment is created, or an existing one is accessed with the `shmget` system call.

# Shared memory (cont.)

- `Char *shmat (int shmid, char *shmaddr, int shmflg);`
- The `shmget` does not provide access to the segment for the calling process.
- We must attach the shared memory segment by calling the `shmat` system call.



# Shared memory (cont.)

- `int shmdt (char *shmaddr);`
- When a process is finished with a shared memory segment, it detaches the segment by calling the `shmdt` system call.
- This call does not delete the shared memory segment.

# Shared memory (cont.)

- `int shmctl (int shmid, int cmd, struct shmid_ds *buf);`
- The `msgctl` function performs various operations in a shared memory segment.

# Inter Process Communication (IPC)

## mapped memory




# Inter Process Communication (IPC)

## pipe



# pipe

- It provides a one-way flow of data.
- It is in the kernel
- It can only be used between processes that have a **parent** process in common.



# pipe (cont.)

- `int pipe (int *filedes);`
- `filedes[0]` : open for reading
- `filedes[1]` : open for writing
- pipe command :
  - `who | sort | lpr`

# Inter Process Communication (IPC)

## fifo



# FIFO

- It is similar to a pipe.
- Unlike pipes, a FIFO has a name associated with it ([named pipe](#)).
- It uses a file as a communication way.
- `int mknod (char *path, int mode, int dev)`
  - mode is or'ed with `S_IFIFO`
  - dev is equal 0 for FIFO.



# Inter Process Communication (IPC)

## socket



# Inter Process Communication (IPC)

## Message passing



# Message Queue

- Message queues are a **linked list** of messages stored within the kernel.
- We don't have to fetch messages in a first-in, first-out order.
  - We can fetch messages based on their type field.
- A process wants to impose some structure on the data being transferred.

# Message Queue (cont.)

- `int msgget (key_t key, int msgflag);`
- A new queue is created, or an existing queue is open by `msgget`.

# Message Queue (cont.)

- `int msgsnd(int msgid, void *ptr, size_t len, int flag);`
- Data is placed onto a message queue by calling `msgsnd`;
- `ptr` points to a long integer that contains the positive integer message `type`, and it is immediately followed by the message `data`.

- Struct `my_msg`

```
{  
    long type;  
    char data[SIZE];  
}
```

# Message Queue (cont.)

- `int msgrcv (int msgid, void *ptr, size_t len, long mtype, int flag);`
- The type argument lets us specify which message we want:
  - `mtype == 0`, the first message on the queue
  - `mtype > 0`, the first message on the queue whose type equals `mtype`.
  - `mtype < 0`, the first message on the queue whose type is the lowest value less or equal to the absolute value of `mtype`.

# Message Queue (cont.)

- `int msgctl (int msgid, int cmd, struct msgid_ds *buf);`
- The `msgctl` function performs various operations in a queue.

# Inter Process Communication (IPC)

## Shared semaphore





# Semaphore

- Semaphores are a synchronization primitive.
- To obtain a shared resource:
  - Test the semaphore that controls the resource.
  - If the value is positive the process can use the resource. The process decrements the value by 1.
  - If the value is 0, the process goes to sleep until the value is greater than 0.

# Semaphore (cont.)

- `int semget (key_t key, int nsems, int flag);`
- This function get a semaphore ID.
- `int semctl (int semid, int semnum, int cmd, union semun arg);`
- The `semctl` function performs various operations in a semaphore.

# Semaphore (cont.)

- `int semop (int semid, struct sembuf *semop, size_t nops);`
- Struct `sembuf`
  - {
    - `ushort sem_num;`
    - `short sem_op;`
    - `shoet sem_flag;`
  - }

# Semaphore (cont.)

- Each particular operation is specified by a `sem_op` value:
  - `sem_op > 0`, this correspond to the release of resources. The `sem_op` value is added to the current value;
  - `sem_op == 0`, the caller wants to wait until the semaphore's value become zero.
  - `sem_op < 0`, this correspond to the allocation of resources. The caller wants to wait until the value become greater or equal the absolute value of `sem_op`. then the absolute value of `sem_op` is subtracted from the current value.